

ORPO: Odds Ratio Preference Optimization

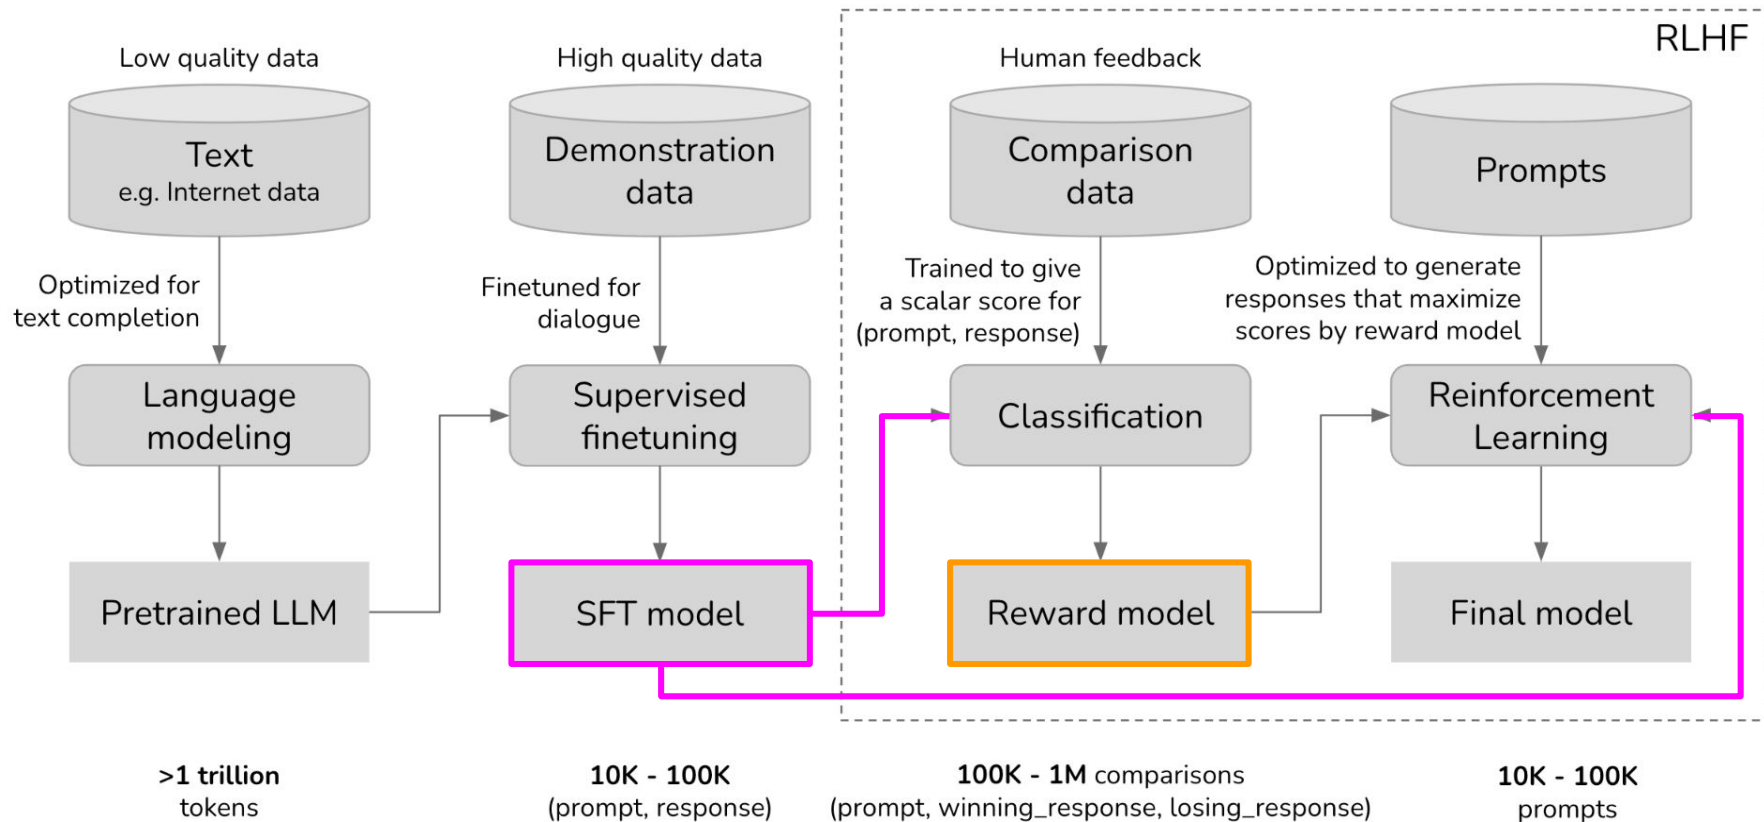
Monolithic Preference Optimization without Reference Model

Outline

- Main takeaways
- SFT without ORPO
- ORPO Loss
- Gradient of ORPO Loss
- Training Details
- Results
- Limitations
- Related Work
- Future Work
- Extra Slides

Main Takeaways

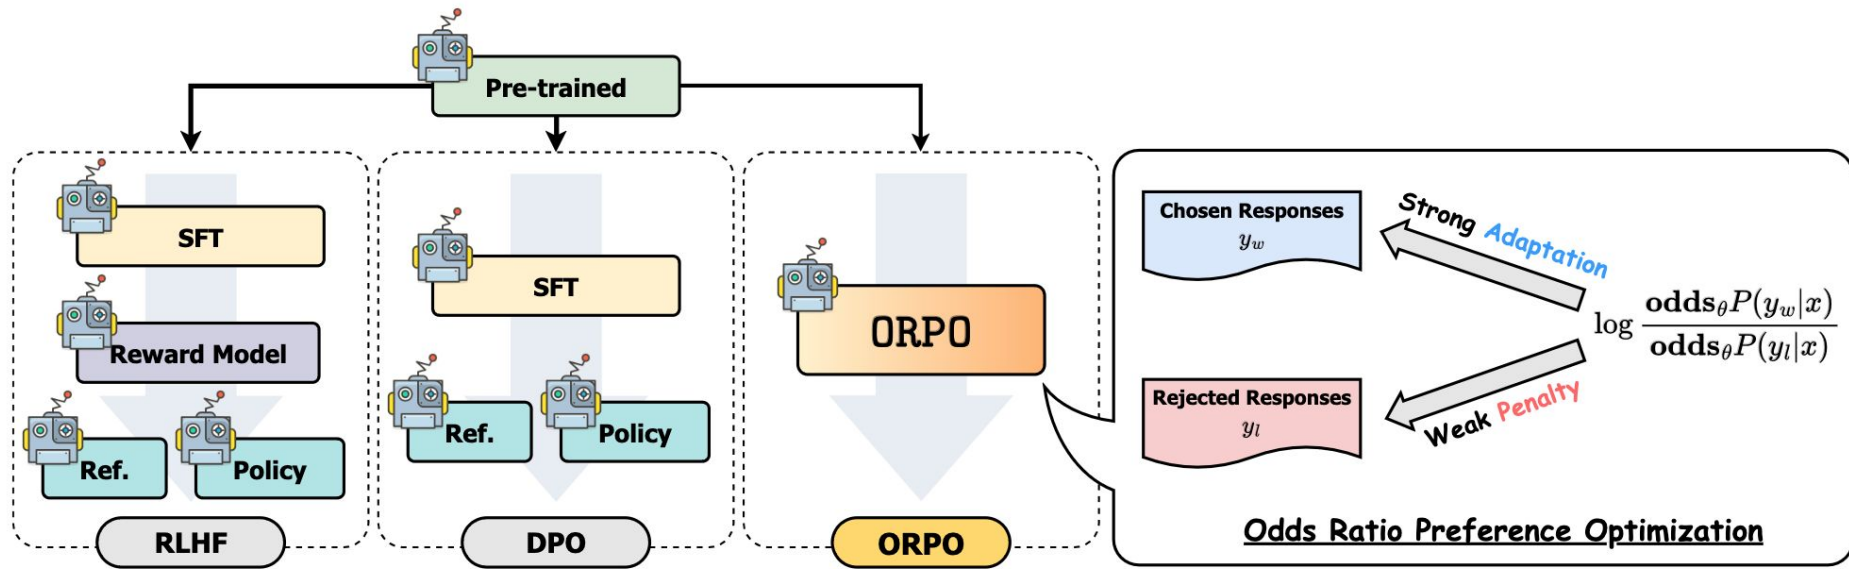
- ORPO does not use a reference model like DPO (KL term) or a reward model and initial SFT model like RLHF.
- ORPO produced a preference-aligned SFT directly.
- The ORPO loss includes a penalty (added to the normal causal LM NLL loss) which maximizes the likelihood of generating a favored response.
- ORPO consistently preferred by a reward model against SFT and RLHF
- ORPO win rate vs. DPO increases as model size increases.



Examples
Bolded: open sourced

Model	Examples
Pretrained LLM	GPT-x, Gopher, Falcon , LLaMa, Pythia , Bloom , StableLM
SFT model	Dolly-v2 , Falcon-Instruct
Reinforcement Learning / Final model	InstructGPT, ChatGPT, Claude, StableVicuna

$$\mathcal{L}_{\text{DPO}}(\pi_{\theta}; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l) \sim \mathcal{D}} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w \mid x)}{\pi_{\text{ref}}(y_w \mid x)} - \beta \log \frac{\pi_{\theta}(y_l \mid x)}{\pi_{\text{ref}}(y_l \mid x)} \right) \right]$$



Using the ORPOTrainer

For a detailed example have a look at the examples/scripts/orpo.py script. At a high level we need to initialize the ORPOTrainer with a model we wish to train. **Note that ORPOTrainer eliminates the need to use the reference model, simplifying the optimization process.** The beta refers to the hyperparameter λ in eq. (6) of the paper and refers to the weighting of the relative odd ratio loss in the standard cross-entropy loss used for SET

```
orpo_config = ORPOConfig(
    beta=0.1, # the lambda/alpha hyperparameter in the paper/code
)

orpo_trainer = ORPOTrainer(
    model,
    args=orpo_config,
    train_dataset=train_dataset,
    tokenizer=tokenizer,
)
```

After this one can then call:

```
orpo_trainer.train()
```

```
orpo_dataset_dict = {
    "prompt": [
        "hello",
        "how are you",
        "What is your name?",
        "What is your name?",
        "Which is the best programming language?",
        "Which is the best programming language?",
        "Which is the best programming language?",
    ],
    "chosen": [
        "hi nice to meet you",
        "I am fine",
        "My name is Mary",
        "My name is Mary",
        "Python",
        "Python",
        "Java",
    ],
    "rejected": [
        "leave me alone",
        "I am not fine",
        "Whats it to you?",
        "I dont have a name",
        "Javascript",
        "C++",
        "C++",
    ],
}
```

SFT without ORPO

$$\mathcal{L} = -\frac{1}{m} \sum_{k=1}^m \log P(\mathbf{x}^{(k)}, \mathbf{y}^{(k)}) \quad (1)$$

$$= -\frac{1}{m} \sum_{k=1}^m \sum_{i=1}^{|V|} y_i^{(k)} \cdot \log(p_i^{(k)}) \quad (2)$$

Effective for domain adaptation
but doesn't have a mechanism to
penalized rejected responses.

y_i is a boolean value that indicates if the i th token in the vocabulary set V is a label token

p_i refers to the probability of the i th token

m is the length of sequence



Figure 3: Log probabilities for chosen and rejected responses during OPT-350M model fine-tuning on HH-RLHF dataset. Despite only chosen responses being used for supervision, rejected responses show a comparable likelihood of generation.

ORPO Loss

$$\mathcal{L}_{ORPO} = \mathbb{E}_{(x, y_w, y_l)} [\mathcal{L}_{SFT} + \lambda \cdot \mathcal{L}_{OR}] \quad (6)$$

Adapt to the specified
subset of the desired
domain

Disfavor generations in the
rejected responses set

$$\mathcal{L}_{OR} = -\log \sigma \left(\log \frac{\mathbf{odds}_{\theta}(y_w|x)}{\mathbf{odds}_{\theta}(y_l|x)} \right) \quad (7)$$

ORPO Loss - TRL

```
592 ~ def odds_ratio_loss(
593     self,
594     policy_chosen_logps: torch.FloatTensor,
595     policy_rejected_logps: torch.FloatTensor,
596 ) -> Tuple[torch.FloatTensor, torch.FloatTensor, torch.FloatTensor, torch.FloatTensor, torch.FloatTensor]:
597     """Compute ORPO's odds ratio (OR) loss for a batch of policy and reference model log probabilities.
598
599     Args:
600         policy_chosen_logps: Log probabilities of the policy model for the chosen responses. Shape: (batch_size,)
601         policy_rejected_logps: Log probabilities of the policy model for the rejected responses. Shape: (batch_size,)
602
603     Returns:
604         A tuple of three tensors: (losses, chosen_rewards, rejected_rewards).
605         The losses tensor contains the ORPO loss for each example in the batch.
606         The chosen_rewards and rejected_rewards tensors contain the rewards for the chosen and rejected responses, respectively.
607         The log odds ratio of the chosen responses over the rejected responses ratio for logging purposes.
608         The `log(sigmoid(log_odds_chosen))` for logging purposes.
609     """
610
611     # Derived from Eqs. (4) and (7) from https://arxiv.org/abs/2403.07691 by using log identities and  $\exp(\log(P(y|x))) = P(y|x)$ 
612     log_odds = (policy_chosen_logps - policy_rejected_logps) - (
613         torch.log1p(-torch.exp(policy_chosen_logps)) - torch.log1p(-torch.exp(policy_rejected_logps))
614     )
615     sig_ratio = F.sigmoid(log_odds)
616     ratio = torch.log(sig_ratio)
617     losses = self.beta * ratio
618
619     chosen_rewards = self.beta * (policy_chosen_logps.to(self.accelerator.device)).detach()
620     rejected_rewards = self.beta * (policy_rejected_logps.to(self.accelerator.device)).detach()
621
622     return losses, chosen_rewards, rejected_rewards, torch.mean(ratio).item(), torch.mean(log_odds).item()
```

ORPO Loss - TRL

```
694 ✓ def cross_entropy_loss(logits, labels):
695     if not self.is_encoder_decoder:
696         # Shift so that tokens < n predict n
697         logits = logits[..., :-1, :].contiguous()
698         labels = labels[..., 1:].contiguous()
699         # Flatten the tokens
700         loss_fct = nn.CrossEntropyLoss()
701         logits = logits.view(-1, logits.shape[-1])
702         labels = labels.view(-1)
703         # Enable model parallelism
704         labels = labels.to(logits.device)
705         loss = loss_fct(logits, labels)
706         return loss
707
708     if self.is_encoder_decoder:
709         labels = concatenated_batch["concatenated_labels"].clone()
710     else:
711         labels = concatenated_batch["concatenated_input_ids"].clone()
712
713     chosen_nll_loss = cross_entropy_loss(all_logits[:len_chosen], labels[:len_chosen])
714
715     all_logps = self.get_batch_logps(
716         all_logits,
717         concatenated_batch["concatenated_labels"],
718         average_log_prob=True,
719         is_encoder_decoder=self.is_encoder_decoder,
720         label_pad_token_id=self.label_pad_token_id,
721     )
722
723     chosen_logps = all_logps[:len_chosen]
724     rejected_logps = all_logps[len_chosen:]
725
726     chosen_logits = all_logits[:len_chosen]
727     rejected_logits = all_logits[len_chosen:]
728
729     return (chosen_logps, rejected_logps, chosen_logits, rejected_logits, chosen_nll_loss)
```

ORPO Loss - TRL

```
730
731 ✓ def get_batch_loss_metrics(
732     self,
733     model,
734     batch: Dict[str, Union[List, torch.LongTensor]],
735     train_eval: Literal["train", "eval"] = "train",
736 ):
737     """Compute the ORPO loss and other metrics for the given batch of inputs for train or test."""
738     metrics = {}
739
740     (
741         policy_chosen_logps,
742         policy_rejected_logps,
743         policy_chosen_logits,
744         policy_rejected_logits,
745         policy_nll_loss,
746     ) = self.concatenated_forward(model, batch)
747
748     losses, chosen_rewards, rejected_rewards, log_odds_ratio, log_odds_chosen = self.odds_ratio_loss(
749         policy_chosen_logps, policy_rejected_logps
750     )
751     # full ORPO loss
752     loss = policy_nll_loss - losses.mean()
753
754     reward accuracies = (chosen_rewards > rejected_rewards).float()
755
756     prefix = "eval_" if train_eval == "eval" else ""
757     metrics[f"{prefix}rewards/chosen"] = chosen_rewards.mean().cpu()
758     metrics[f"{prefix}rewards/rejected"] = rejected_rewards.mean().cpu()
759     metrics[f"{prefix}rewards/accuracies"] = reward accuracies.mean().cpu()
760     metrics[f"{prefix}rewards/margins"] = (chosen_rewards - rejected_rewards).mean().cpu()
761     metrics[f"{prefix}logps/rejected"] = policy_rejected_logps.detach().mean().cpu()
762     metrics[f"{prefix}logps/chosen"] = policy_chosen_logps.detach().mean().cpu()
763     metrics[f"{prefix}logits/rejected"] = policy_rejected_logits.detach().mean().cpu()
764     metrics[f"{prefix}logits/chosen"] = policy_chosen_logits.detach().mean().cpu()
765     metrics[f"{prefix}nll_loss"] = policy_nll_loss.detach().mean().cpu()
766     metrics[f"{prefix}log_odds_ratio"] = log_odds_ratio
767     metrics[f"{prefix}log_odds_chosen"] = log_odds_chosen
768
```

The odds of generating the output sequence y given an input sequence x :

$$\mathbf{odds}_\theta(y|x) = \frac{P_\theta(y|x)}{1 - P_\theta(y|x)}$$

Intuitively $\mathbf{odds}_\theta(y|x) = k$ implies that it is k times more likely for the model θ to generate the output sequence y than not generating it.

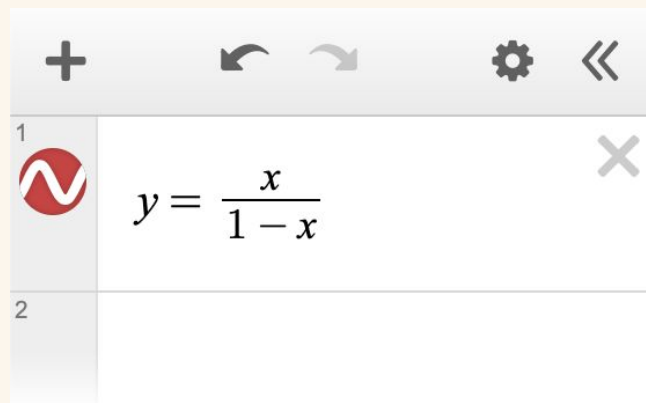
Writing that out:

$$\mathbf{odds}_\theta(y|x) = \frac{P_\theta(y|x)}{1 - P_\theta(y|x)} = k$$


$$\frac{P_\theta(y|x)}{1 - P_\theta(y|x)} = k$$

$$P_\theta(y|x) = k[1 - P_\theta(y|x)]$$


it is k times more likely to generate y (i.e. P_θ) than not generating it ($1 - P_\theta$).



$$\mathbf{odds}_{\theta}(y|x) = \frac{P_{\theta}(y|x)}{1 - P_{\theta}(y|x)}$$

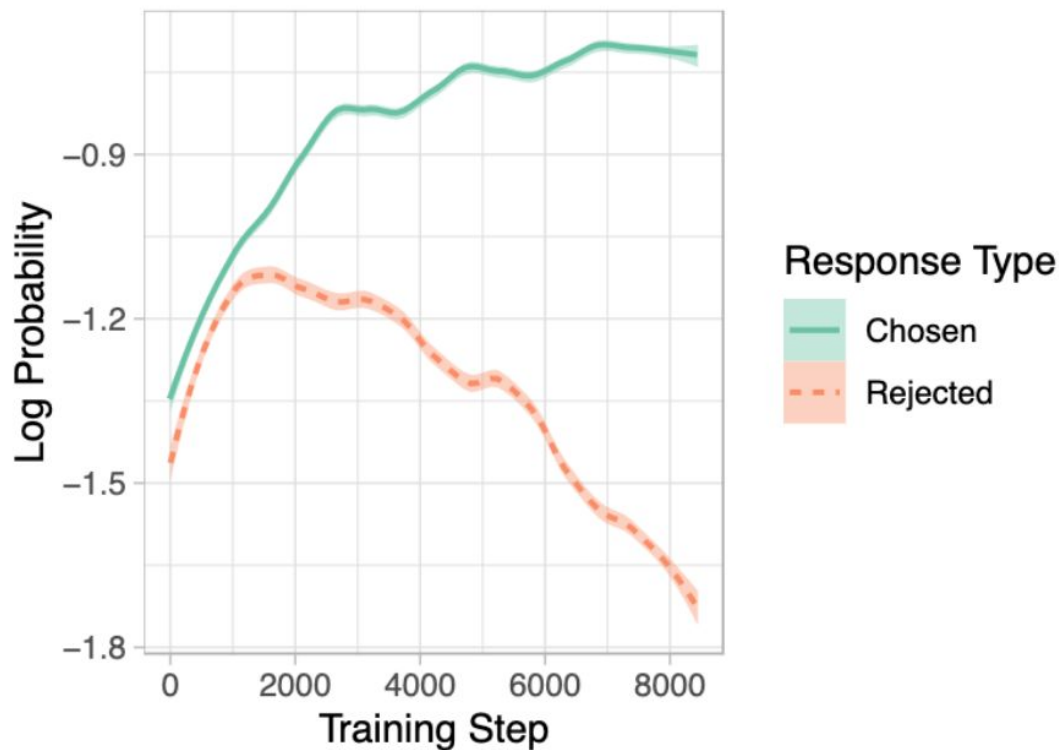



$$\mathcal{L}_{OR} = -\log \sigma \left(\log \frac{\mathbf{odds}_{\theta}(y_w|x)}{\mathbf{odds}_{\theta}(y_l|x)} \right) \quad (7)$$

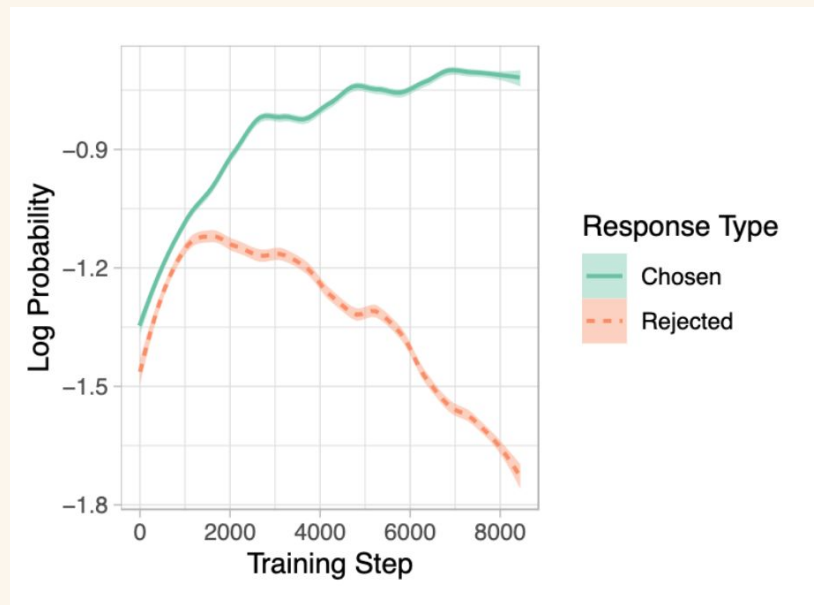


Minimizing \mathcal{L}_{OR} means maximizing $\frac{\mathbf{odds}_{\theta}(y_w|x)}{\mathbf{odds}_{\theta}(y_l|x)}$

$$\mathcal{L}_{OR} = -\log \sigma \left(\log \frac{\mathbf{odds}_{\theta}(y_w|x)}{\mathbf{odds}_{\theta}(y_l|x)} \right) \quad (7)$$



$$\mathcal{L}_{ORPO} = \mathbb{E}_{(x, y_w, y_l)} [\mathcal{L}_{SFT} + \lambda \cdot \mathcal{L}_{OR}] \quad (6)$$



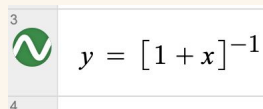
Gradient of ORPO Loss

$$\nabla_{\theta} \mathcal{L}_{OR} = \delta(d) \cdot h(d) \quad (8)$$


$$\delta(d) = \left[1 + \frac{\mathbf{odds}_{\theta} P(y_w|x)}{\mathbf{odds}_{\theta} P(y_l|x)} \right]^{-1} \quad (9)$$

$$h(d) = \frac{\nabla_{\theta} \log P_{\theta}(y_w|x)}{1 - P_{\theta}(y_w|x)} - \frac{\nabla_{\theta} \log P_{\theta}(y_l|x)}{1 - P_{\theta}(y_l|x)} \quad (10)$$

Gradient of ORPO Loss: $\delta(d)$



$y = [1 + x]^{-1}$



$$\delta(d) = \left[1 + \frac{\mathbf{odds}_{\theta} P(y_w | x)}{\mathbf{odds}_{\theta} P(y_l | x)} \right]^{-1}$$

When the odds of the favored response are relatively higher than the disfavored responses, $\delta(d)$ will converge to 0.

Gradient of ORPO Loss: $\delta(d)$

$$\delta(d) = \left[1 + \frac{\mathbf{odds}_{\theta} P(y_w|x)}{\mathbf{odds}_{\theta} P(y_l|x)} \right]^{-1}$$

$$\delta(d) = \left[\frac{\mathbf{odds}_{\theta} P(y_l|x)}{\mathbf{odds}_{\theta} P(y_l|x)} + \frac{\mathbf{odds}_{\theta} P(y_w|x)}{\mathbf{odds}_{\theta} P(y_l|x)} \right]^{-1}$$

$$\delta(d) = \left[\frac{\mathbf{odds}_{\theta} P(y_l|x) + \mathbf{odds}_{\theta} P(y_w|x)}{\mathbf{odds}_{\theta} P(y_l|x)} \right]^{-1}$$

$$\delta(d) = \left[\frac{\mathbf{odds}_{\theta} P(y_l|x)}{\mathbf{odds}_{\theta} P(y_l|x) + \mathbf{odds}_{\theta} P(y_w|x)} \right]$$

$\delta(d)$ will play the role of a penalty term, accelerating the parameter updates if the model is more likely to generate the rejected responses.

Gradient of ORPO Loss: $h(d)$

$$h(d) = \frac{\nabla_{\theta} \log P_{\theta}(y_w|x)}{1 - P_{\theta}(y_w|x)} - \frac{\nabla_{\theta} \log P_{\theta}(y_l|x)}{1 - P_{\theta}(y_l|x)} \quad (10)$$

For the chosen responses, this accelerates the model's adaptation toward the distribution of chosen responses as the likelihood increases.

Gradient of ORPO Loss

$$\nabla_{\theta} \mathcal{L}_{OR} = \delta(d) \cdot h(d) \quad (8)$$

$$\delta(d) = \left[1 + \frac{\mathbf{odds}_{\theta} P(y_w|x)}{\mathbf{odds}_{\theta} P(y_l|x)} \right]^{-1}$$

accelerates the parameter updates if the model is more likely to generate the rejected responses

$$h(d) = \frac{\nabla_{\theta} \log P_{\theta}(y_w|x)}{1 - P_{\theta}(y_w|x)} - \frac{\nabla_{\theta} \log P_{\theta}(y_l|x)}{1 - P_{\theta}(y_l|x)}$$

accelerates the model's adaptation toward the distribution of chosen responses as the likelihood increases

(10)

Gradient of ORPO Loss

$$\nabla_{\theta} \mathcal{L}_{OR} = \delta(d) \cdot h(d) \quad (8)$$

$$\delta(d) = \left[1 + \frac{\mathbf{odds}_{\theta} P(y_w|x)}{\mathbf{odds}_{\theta} P(y_l|x)} \right]^{-1}$$

larger when rejected odds are larger than favored odds

$$h(d) = \frac{\nabla_{\theta} \log P_{\theta}(y_w|x)}{1 - P_{\theta}(y_w|x)} - \frac{\nabla_{\theta} \log P_{\theta}(y_l|x)}{1 - P_{\theta}(y_l|x)} \quad (10)$$

larger when favored odds are larger than rejected odds

Training - Models

- OPT (125M to 1.3B) for SFT, PPO, DPO and ORPO
- Phi-2 (2.7B)
- Llama 2 (7B)
- Mistral (7B)

Training - Hyperparameters

- Flash Attention 2
- OPT series and Phi-2: Deep Speed Zero 2
- Llama-2 (7B) and Mistral (7B): FSDP
- Optimizer: AdamW and paged Adam
- Linear Warmup with Cosine Decay
- Input Length: 1024 (HH), 2048 (UltraFeedback)

Training - Hyperparameters

- SFT: Max LR = $1e-5$, 1 epoch
- DPO: $\beta = 0.1$, LR= $5e-6$, 3 epochs
- ORPO: LR= $8e-6$, 10 epochs

Hyperparameter	Setting
ppo_epoch	4
init_kl_coef	0.1
horizon	2,000
batch_size	64
mini_batch_size	8
gradient_accumulation_steps	1
output_min_length	128
output_max_length	512
optimizer	AdamW
learning_rate	$1e-05$
gamma	0.99

Table 5: Hyperparameter settings for RLHF.

Training - Datasets

- Anthropic's HH-RLHF
- Binarized UltraFeedback

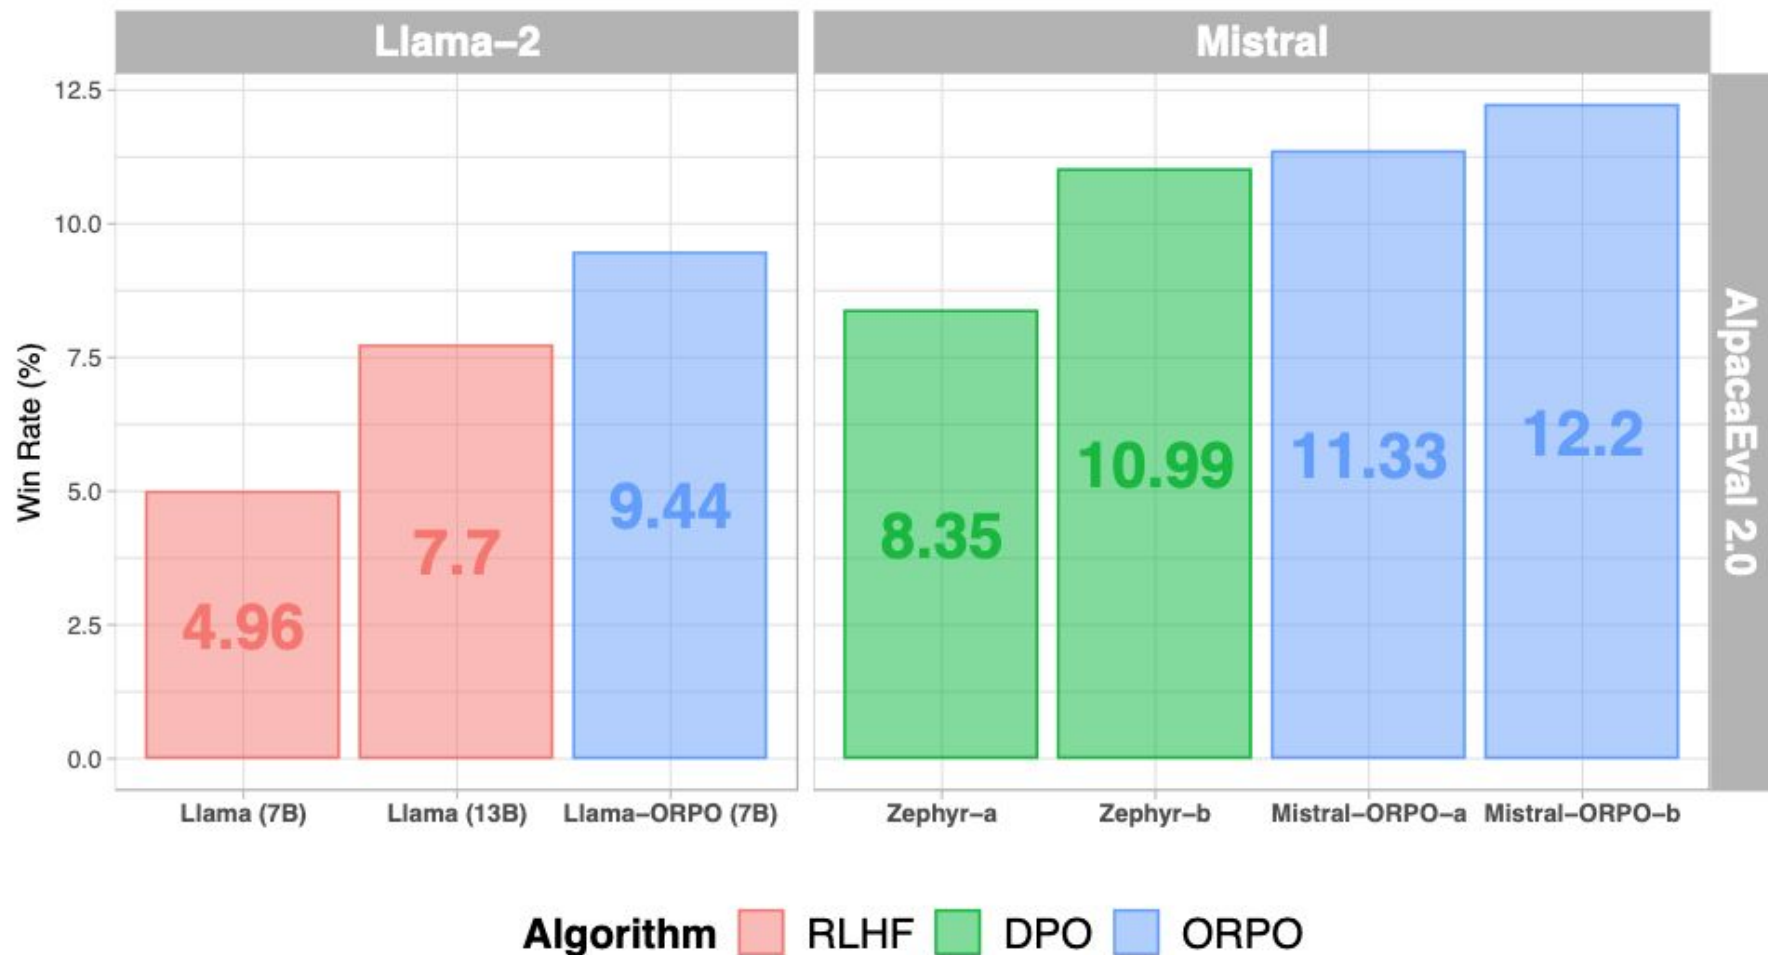
Training - Reward Model

- OPT 250M and 1.3B on each dataset for a single epoch


$$-\mathbb{E}_{(x, y_l, y_w)} [\log \sigma (r(x, y_w) - r(x, y_l))] \quad (11)$$

Results: AlpacaEval

Model Name	Size	AlpacaEval _{1.0}	AlpacaEval _{2.0}
Phi-2 + SFT	2.7B	48.37% (1.77)	0.11% (0.06)
Phi-2 + SFT + DPO	2.7B	50.63% (1.77)	0.78% (0.22)
Phi-2 + ORPO (<i>Ours</i>)	2.7B	71.80% (1.59)	6.35% (0.74)
Llama-2 Chat *	7B	71.34% (1.59)	4.96% (0.67)
Llama-2 Chat *	13B	81.09% (1.38)	7.70% (0.83)
Llama-2 + ORPO (<i>Ours</i>)	7B	81.26% (1.37)	9.44% (0.85)
Zephyr (α) *	7B	85.76% (1.23)	8.35% (0.87)
Zephyr (β) *	7B	90.60% (1.03)	10.99% (0.96)
Mistral-ORPO- α (<i>Ours</i>)	7B	87.92% (1.14)	11.33% (0.97)
Mistral-ORPO- β (<i>Ours</i>)	7B	91.41% (1.15)	12.20% (0.98)



Results: MT-Bench



Results: MT-Bench (varying λ)




Figure 10: MT-Bench result comparison by differing $\lambda = 0.1$ and $\lambda = 1.0$.

Results: MT-Bench (varying λ)

$$\mathcal{L}_{ORPO} = \mathbb{E}_{(x, y_w, y_l)} [\mathcal{L}_{SFT} + \lambda \cdot \mathcal{L}_{OR}]$$




Figure 10: MT-Bench result comparison by differing $\lambda = 0.1$ and $\lambda = 1.0$.

In comparison to $\lambda = 0.1$, Mistral+ORPO (7B) with $\lambda = 1.0$ performs worse in extraction, math, and reasoning, which are the categories that generally require deterministic answers. On the other hand, it performs better in STEM, humanities, and roleplay, which ask the generations without hard answers.

Results: HH-RLHF

ORPO vs	SFT	+DPO	+PPO
OPT-125M	84.0 (0.62)	41.7 (0.77)	66.1 (0.26)
OPT-350M	82.7 (0.56)	49.4 (0.54)	79.4 (0.29)
OPT-1.3B	78.0 (0.16)	70.9 (0.52)	65.9 (0.33)

Table 2: Average win rate (%) and its standard deviation of ORPO and standard deviation over other methods on **HH-RLHF** dataset for three rounds. Sampling decoding with a temperature of 1.0 was used on the test set.

Results: UltraFeedback

ORPO vs	SFT	+DPO	+PPO
OPT-125M	73.2 (0.12)	48.8 (0.29)	71.4 (0.28)
OPT-350M	80.5 (0.54)	50.5 (0.17)	85.8 (0.62)
OPT-1.3B	69.4 (0.57)	57.8 (0.73)	65.7 (1.07)

Table 3: Average win rate (%) and its standard deviation of ORPO and standard deviation over other methods on **UltraFeedback** dataset for three rounds. Sampling decoding with a temperature of 1.0 was used.

Results: Reward Distribution




Figure 5: Reward distribution comparison between OPT-125M (left), OPT-350M (middle), and OPT-1.3B (right) trained with SFT (blue), RLHF (green), DPO (orange), and ORPO (red) on the test set of UltraFeedback using the RM-1.3B. While the rewards of the trained models are roughly normal and preference optimization algorithms (RLHF, DPO, and ORPO) tend to move the reward distribution in the positive direction, ORPO is on par or better than RLHF and DPO in increasing the expected reward. The same plot for the HH-RLHF dataset is in Appendix F.

Results: Lexical Diversity

	Per Input↓	Across Input↓
Phi-2 + SFT + DPO	0.8012	0.6019
Phi-2 + ORPO	0.8909	0.5173
Llama-2 + SFT + DPO	0.8889	0.5658
Llama-2 + ORPO	0.9008	0.5091

Table 4: Lexical diversity of Phi-2 and Llama-2 fine-tuned with DPO and ORPO. Lower cosine similarity is equivalent to higher diversity. The highest value in each column within the same model family is bolded.

Results: Computational Efficiency

- ORPO does not require a reference model which in RLHF and DPO is the model trained with SFT used as the baseline model for updating the parameters.
- DPO and RLHF require two SFT models: a frozen reference model (KL term) and the model being trained.
- ORPO only has one model: the model being trained with SFT. This requires half the forward passes of DPO or RLHF.

Limitations

- Lacks comparisons with alignment algorithms other than PPO and DPO
- Only trained models up to 7B parameters

Future Work

- Evaluate performance on models larger than 7B parameters
- Evaluate impact of ORPO on pretrained model
- Expand to consecutive preference alignment algorithms

Related Work

- Alignment with Reinforcement Learning (RLHF)
- Alignment without Reward Model (DPO)
- Alignment with Supervised Fine-Tuning (filtered datasets)

Extra Slides

Why not Probability Ratio?

$$\mathbf{PR}_{\theta}(y_w, y_l) = \frac{P_{\theta}(y_w|x)}{P_{\theta}(y_l|x)} \quad (16)$$

The probability ratio leads to more extreme discrimination of the disfavored responses than the odds ratio.

The excessive margin could lead to the unwarranted suppression of logits for tokens in disfavored responses within the incorporated setting, potentially resulting in issues of degeneration.




Figure 6: Sampled distribution of $\log \mathbf{PR}(X_2|X_1)$ and $\log \mathbf{OR}(X_2|X_1)$. $\log \mathbf{OR}(X_2|X_1)$ has a wider range given the same input probability pairs (X_1, X_2) .

Weighting Value λ

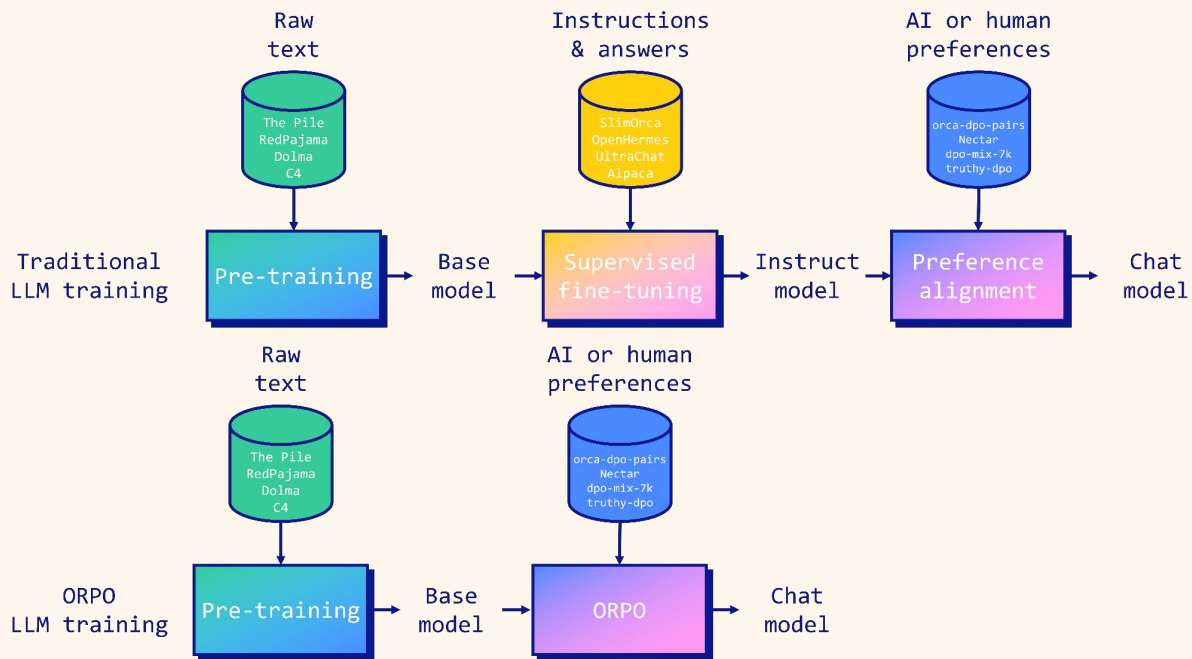
$$\mathcal{L}_{ORPO} = \mathbb{E}_{(x, y_w, y_l)} [\mathcal{L}_{SFT} + \lambda \cdot \mathcal{L}_{OR}]$$



Figure 9: The log probability trend by λ . With larger λ (e.g., $\lambda = 1.0$), \mathcal{L}_{OR} gets more influential in fine-tuning the models with ORPO.

Fine-tuning Llama 3 with ORPO

<https://huggingface.co/blog/mlabonne/orpo-llama-3>



Math Notebook

https://colab.research.google.com/drive/1dgR8O4pbuvecVEkfq6xf_dHwVn2G9Kjf?usp=sharing